

# CAP-372 / 2019 - Quinta Lista de Exercícios

Data: 19/09/19

Este documento e os fontes estão em: <https://github.com/xxxxxxxm/372>

## SUMÁRIO

ALGORITMO.....	1
SIMULAÇÃO COM 4 PROCESSADORES.....	2
SIMULAÇÃO COM 2 PROCESSADORES.....	5
OBTENÇÃO DO TEMPO SERIAL.....	10
PARALELIZANDO.....	11
RODANDO NO SANTOS DUMONT.....	16
TEMPORIZAÇÃO DO PROCESSAMENTO REFERENTE AOS CÁLCULOS E ÀS COMUNICAÇÕES.....	17
SPEED UP.....	21
REFERÊNCIAS.....	22

## ALGORITMO

Testando algoritmo 1 de [http://www.lac.inpe.br/~stephan/CAP-372/matrixmult\\_microsoft.pdf](http://www.lac.inpe.br/~stephan/CAP-372/matrixmult_microsoft.pdf) ;  
compilando na máquina local (Intel i7) usando *gfortran*, sem paralelização:

```
! CAP372 - exercise 5 - 2019-09-15
! Simple check of sequential algorithm 1
! Compiling on the local machine using gfortran
! - Each subtask i hold one row of matrix A and one column of matrix B .
! - The result is stored in one element of C .
! - After that, every subtask i, 0≤i<n, transmits its column of matrix B
!   to the subtask with the number (i + 1) mod n .
! - When the number of processors ("p") is less than the number of
!   basic subtasks ("n"), each processor would execute several
!   inner products of matrix A rows and matrix B columns.
! - Each aggregated basic subtask ( = the calculation of one row of C)
!   determines several rows of the result matrix C.
program list05_algorithm
  implicit none
  integer :: mi, mj, mk           ! matrix indexes
  double precision :: t1, t2      ! time elapsed
  double precision :: temp1       ! temporary
  integer :: p                    ! emulate number of processes
  integer :: stripe_i, stripe_qtd, stripe_count, stripe_countA ! stripe
  integer :: iteration, subtask   ! looping
```

```

! C = A * B, matrix initialization with example
integer, parameter :: ms = 4          ! matrix size
double precision, dimension(ms, ms) ::
  A = transpose(reshape([ 5, 2, 6, 1,
                           0, 6, 2, 0,
                           3, 8, 1, 4,
                           1, 8, 5, 6 ], shape(A))), &
  B = transpose(reshape([ 7, 5, 8, 0,
                           1, 8, 2, 6,
                           9, 4, 3, 8,
                           5, 3, 7, 9 ], shape(B))), &
  C = 0

print*, "=== List05 - Matrix multiplication example ==="
call cpu_time(t1)                    ! elapsed time calculation
p = 4                                ! emulate 4 processors
stripe_qtd = ms / p                  ! subtasks per processor
do iteration = 0, ms / stripe_qtd - 1 ! algorithm iteration
  do subtask = 0, ms / stripe_qtd - 1 ! each subtask hold one matrix row
    do stripe_countA = 0, stripe_qtd - 1 ! A count for p < n
      do stripe_count = 0, stripe_qtd - 1 ! count for p < n
        ! stripe column calculation
        stripe_i = stripe_qtd*mod(subtask+p-iteration,p)+stripe_count
        ! matrix A line calculation
        mi = subtask * stripe_qtd + 1 + stripe_countA
        temp1 = 0 ! multiplication: A line x B column
        do mj = 1, ms ! count A column
          mk = stripe_i + 1 ! B column index
          temp1 = temp1 + A(mi, mj) * B(mj, mk) ! C = A x B
          print*, "A =", int(A(mi, mj)), " ", B =", int(B(mj, mk))
        end do
        print*, "iteration =", iteration, " ", subtask =", subtask, &
          " ", C(i, j) = ", int(temp1)
        C(mi, mk) = temp1 ! store C element
      end do
    end do
  end do
end do
print*, "Result Matrix C =" ! show the result
do mi = 1, ms
  print*, int(C(mi, :))
end do
call cpu_time(t2) ! elapsed time calculation
print *, "Elapsed time [s]:", t2 - t1
end program list05_algorithm

```

Multiplicação de teste:

$$\begin{vmatrix} 5 & 2 & 6 & 1 \\ 0 & 6 & 2 & 0 \\ 3 & 8 & 1 & 4 \\ 1 & 8 & 5 & 6 \end{vmatrix} \times \begin{vmatrix} 7 & 5 & 8 & 0 \\ 1 & 8 & 2 & 6 \\ 9 & 4 & 3 & 8 \\ 5 & 3 & 7 & 9 \end{vmatrix} = \begin{vmatrix} 96 & 68 & 69 & 69 \\ 24 & 56 & 18 & 52 \\ 58 & 95 & 71 & 92 \\ 90 & 107 & 81 & 142 \end{vmatrix}$$

## SIMULAÇÃO COM 4 PROCESSADORES

### INTERAÇÃO 1

```

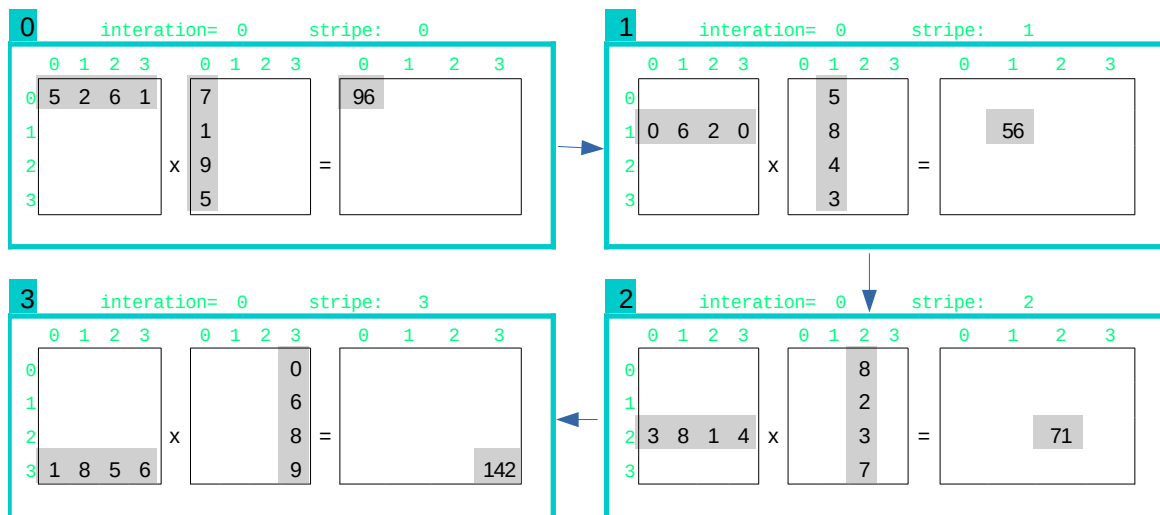
=== List05 - Matrix multiplication example ===
A =          5      , B =          7
A =          2      , B =          1
A =          6      , B =          9
A =          1      , B =          5
iteration =      0      subtask =          0      C(i, j) =
96

```

```

A =      0      , B =      5
A =      6      , B =      8
A =      2      , B =      4
A =      0      , B =      3
iteration =      0      subtask =      1      C(i,j) =
56
A =      3      , B =      8
A =      8      , B =      2
A =      1      , B =      3
A =      4      , B =      7
iteration =      0      subtask =      2      C(i,j) =
71
A =      1      , B =      0
A =      8      , B =      6
A =      5      , B =      8
A =      6      , B =      9
iteration =      0      subtask =      3      C(i,j) =
142
A =      5      , B =      0
A =      2      , B =      6
A =      6      , B =      8
A =      1      , B =      9

```

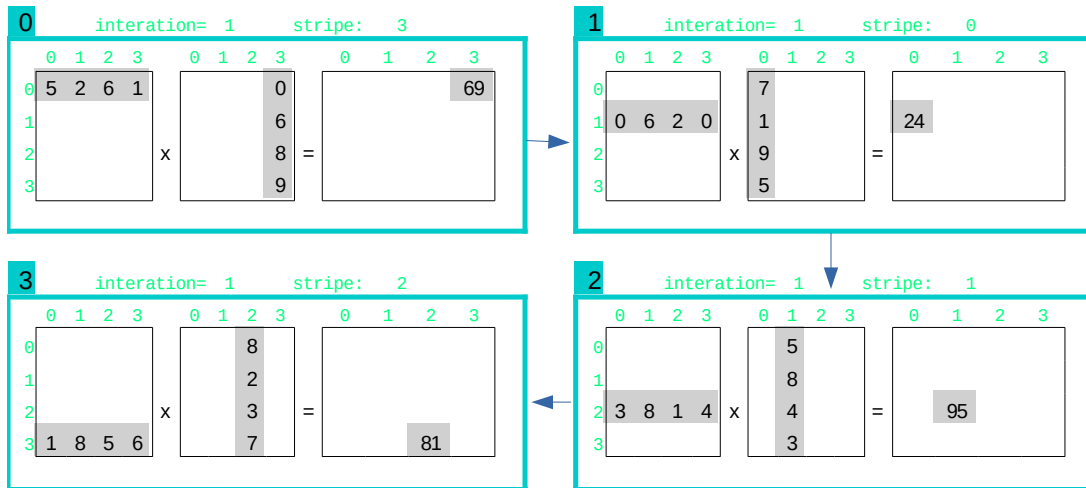


## INTERAÇÃO 2

```

iteration =      1      subtask =      0      C(i,j) =
69
A =      0      , B =      7
A =      6      , B =      1
A =      2      , B =      9
A =      0      , B =      5
iteration =      1      subtask =      1      C(i,j) =
24
A =      3      , B =      5
A =      8      , B =      8
A =      1      , B =      4
A =      4      , B =      3
iteration =      1      subtask =      2      C(i,j) =
95
A =      1      , B =      8
A =      8      , B =      2
A =      5      , B =      3
A =      6      , B =      7
iteration =      1      subtask =      3      C(i,j) =
81
A =      5      , B =      8
A =      2      , B =      2
A =      6      , B =      3
A =      1      , B =      7

```

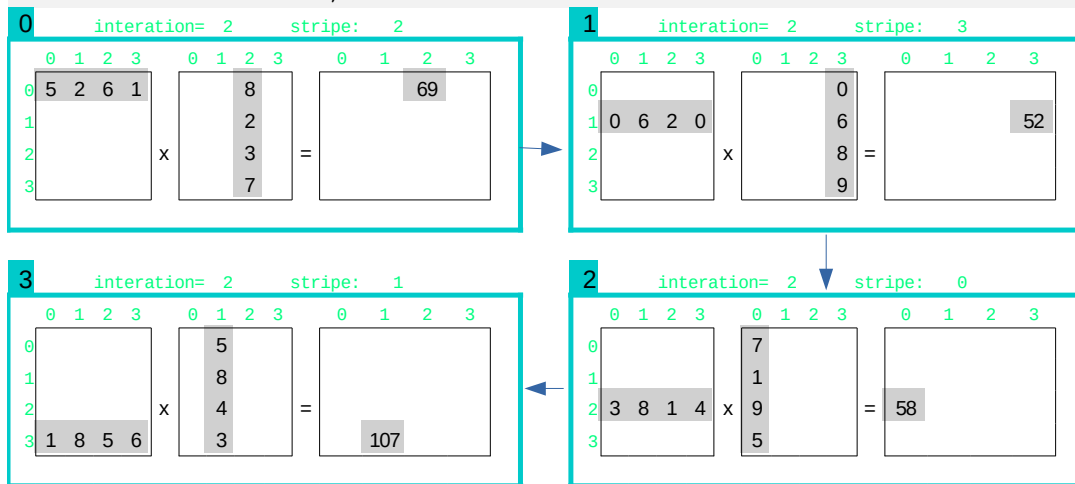


### INTERAÇÃO 3

```

iteration =      2      subtask =      0      C(i,j) =
69
A =      0      , B =      0
A =      6      , B =      6
A =      2      , B =      8
A =      0      , B =      9
iteration =      2      subtask =      1      C(i,j) =
52
A =      3      , B =      7
A =      8      , B =      1
A =      1      , B =      9
A =      4      , B =      5
iteration =      2      subtask =      2      C(i,j) =
58
A =      1      , B =      5
A =      8      , B =      8
A =      5      , B =      4
A =      6      , B =      3
iteration =      2      subtask =      3      C(i,j) =
107
A =      5      , B =      5
A =      2      , B =      8
A =      6      , B =      4
A =      1      , B =      3

```



### INTERAÇÃO 4

```

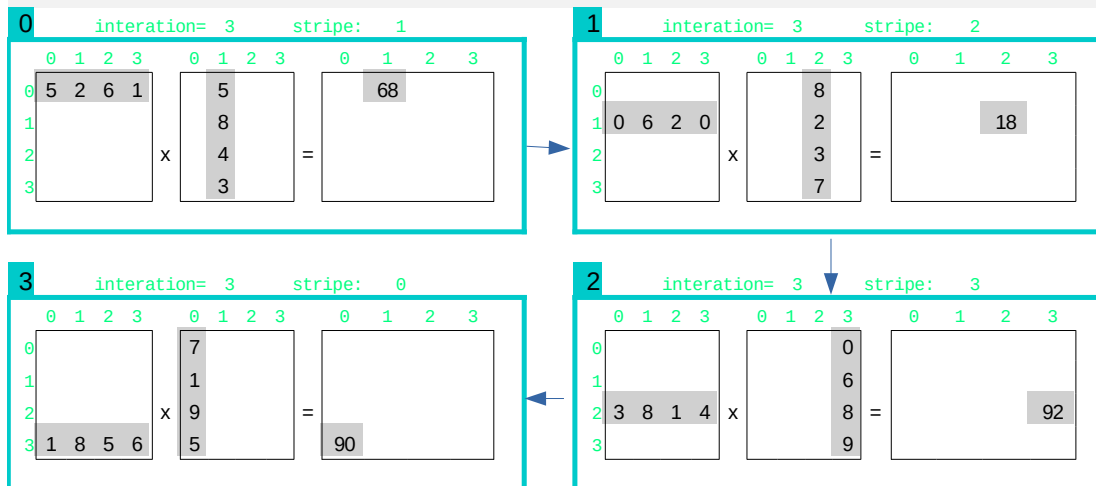
iteration =      3      subtask =      0      C(i,j) =
68
A =      0      , B =      8

```

```

A =          6      , B =          2
A =          2      , B =          3
A =          0      , B =          7
iteration =      3      subtask =          1      C(i,j) =
18
A =          3      , B =          0
A =          8      , B =          6
A =          1      , B =          8
A =          4      , B =          9
iteration =      3      subtask =          2      C(i,j) =
92
A =          1      , B =          7
A =          8      , B =          1
A =          5      , B =          9
A =          6      , B =          5
iteration =      3      subtask =          3      C(i,j) =
90
Result Matrix C =
      96      68      69      69
      24      56      18      52
      58      95      71      92
      90     107      81     142

```



## SIMULAÇÃO COM 2 PROCESSADORES

### Algorithm 1: Block-Striped Decomposition...

#### □ Aggregating and Distributing the Subtasks among the Processors:

- In case when the number of processors  $p$  is less than the number of basic subtasks  $n$ , calculations can be aggregated in such a way that **each processor would execute several inner products** of matrix  $A$  rows and matrix  $B$  columns. In this case after the completion of computation, **each aggregated basic subtask determines several rows of the result matrix  $C$** ,
- Under such conditions the initial matrix  $A$  is decomposed into  $p$  horizontal stripes and matrix  $B$  is decomposed into  $p$  vertical stripes,
- Subtasks distribution among the processors have to meet the requirements of effective representation of the ring structure of subtask information dependencies

### INTERAÇÃO 1

```


=== List05 - Matrix multiplication example ===
A =          5      , B =          7
A =          2      , B =          1
A =          6      , B =          9
A =          1      , B =          5

```

```

iteration =          0          subtask =          0          C(i,j) =
96
A =          5          , B =          5
A =          2          , B =          8
A =          6          , B =          4
A =          1          , B =          3
iteration =          0          subtask =          0          C(i,j) =
68
A =          0          , B =          7
A =          6          , B =          1
A =          2          , B =          9
A =          0          , B =          5
iteration =          0          subtask =          0          C(i,j) =
24
A =          0          , B =          5
A =          6          , B =          8
A =          2          , B =          4
A =          0          , B =          3
iteration =          0          subtask =          0          C(i,j) =
56
A =          3          , B =          8
A =          8          , B =          2
A =          1          , B =          3
A =          4          , B =          7
iteration =          0          subtask =          1          C(i,j) =
71
A =          3          , B =          0
A =          8          , B =          6
A =          1          , B =          8
A =          4          , B =          9
iteration =          0          subtask =          1          C(i,j) =
92
A =          1          , B =          8
A =          8          , B =          2
A =          5          , B =          3
A =          6          , B =          7
iteration =          0          subtask =          1          C(i,j) =
81
A =          1          , B =          0
A =          8          , B =          6
A =          5          , B =          8
A =          6          , B =          9
iteration =          0          subtask =          1          C(i,j) =
142
A =          5          , B =          8
A =          2          , B =          2
A =          6          , B =          3
A =          1          , B =          7
iteration =          1          subtask =          0          C(i,j) =
69
A =          5          , B =          0
A =          2          , B =          6
A =          6          , B =          8
A =          1          , B =          9

```



## INTERAÇÃO 2

```


iteration =          1          subtask =          0          C(i,j) =
69
A =          0          , B =          8
A =          6          , B =          2

```

```

A =      2      , B =      3
A =      0      , B =      7
iteration =    1      subtask =    0      C(i,j) =
18
A =      0      , B =      0
A =      6      , B =      6
A =      2      , B =      8
A =      0      , B =      9
iteration =    1      subtask =    0      C(i,j) =
52
A =      3      , B =      7
A =      8      , B =      1
A =      1      , B =      9
A =      4      , B =      5
iteration =    1      subtask =    1      C(i,j) =
58
A =      3      , B =      5
A =      8      , B =      8
A =      1      , B =      4
A =      4      , B =      3
iteration =    1      subtask =    1      C(i,j) =
95
A =      1      , B =      7
A =      8      , B =      1
A =      5      , B =      9
A =      6      , B =      5
iteration =    1      subtask =    1      C(i,j) =
90
A =      1      , B =      5
A =      8      , B =      8
A =      5      , B =      4
A =      6      , B =      3
iteration =    1      subtask =    1      C(i,j) =
107
Result Matrix C =
      96      68      69      69
      24      56      18      52
      58      95      71      92
      90     107      81     142

```

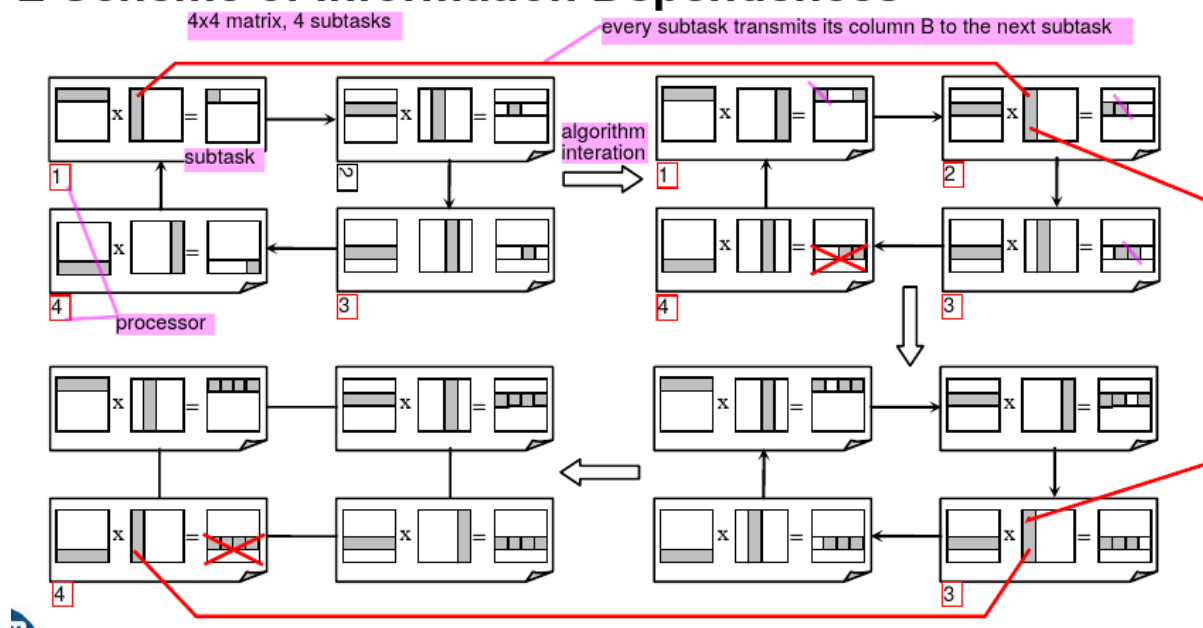






As figuras anteriores derivam do slide:

## ❑ Scheme of Information Dependences



# OBTENÇÃO DO TEMPO SERIAL

Rodando no nó de acesso do Santos Dumont

Programa:

```
program ilist05s512
! CAP372 - exercise 5 - 2019-09-15
! - Each subtask i hold one row of matrix A and one column of matrix B .
! - The result is stored in one element of C .
! - After that, every subtask i, 0≤i<n, transmits its column of matrix B
!   to the subtask with the number (i + 1) mod n .
! - When the number of processors ("p") is less than the number of
!   basic subtasks ("n"), each processor would execute several
!   inner products of matrix A rows and matrix B columns.
! - Each aggregated subtask (= the calculation of one row of C)
!   determines several rows of the result matrix C.
implicit none
integer, parameter :: ms = 512      ! matrix size
integer, parameter :: p = 8        ! emulate the number of processes
integer :: mi, mj, mk               ! matrix indexes
double precision :: t1, t2          ! time elapsed
double precision :: temp1           ! temporary
integer :: stripe_i, stripe_qtd, stripe_count, stripe_countA ! stripe
integer :: iteration, subtask       ! looping

! C = A * B, matrix initialization
double precision, dimension(ms, ms) :: C = 0, &
A = reshape([(mi, mi=1, ms*ms)], shape(A)), &
B = reshape([(mi, mi=ms*ms+1, ms*ms*2)], shape(A))

call cpu_time(t1)                  ! elapsed time calculation
stripe_qtd = ms / p                ! subtasks per processor
do iteration = 0, ms / stripe_qtd - 1 ! algorithm iteration
  do subtask = 0, ms / stripe_qtd - 1 ! each subtask hold one matrix row
    do stripe_countA = 0, stripe_qtd - 1 ! A count for p < n
      do stripe_count = 0, stripe_qtd - 1 ! count for p < n
        ! stripe column calculation
        stripe_i = stripe_qtd*mod(subtask+p-iteration,p)+stripe_countA
        ! matrix A line calculation
        mi = subtask * stripe_qtd + 1 + stripe_countA
        temp1 = 0 ! multiplication: A line x B column
        do mj = 1, ms ! count A column
          mk = stripe_i + 1 ! B column index
          temp1 = temp1 + A(mi, mj) * B(mj, mk) ! C = A x B
        end do
        C(mi, mk) = temp1 ! store C element
      end do
    end do
  end do
end do
call cpu_time(t2)                  ! elapsed time calculation
write(*, "(A, I5, A, I3, A, F8.4)") 'Serial version: matrix size:', ms, &
', Emulated Processes:', p, ', Elapsed time [s]:', t2 - t1
end program ilist05s512
```

Resultados:

```
Serial version: matrix size: 512, Emulated Processes: 2, Elapsed time [s]:
1.3891
[xxxxxxx.xxxxxxx@s dumont13 ~]$ ./ilist05s512p4
Serial version: matrix size: 512, Emulated Processes: 4, Elapsed time [s]:
1.3570
[xxxxxxx.xxxxxxx@s dumont13 ~]$ ./ilist05s512p8
Serial version: matrix size: 512, Emulated Processes: 8, Elapsed time [s]:
1.3226
```

# PARALELIZANDO

## Programa:

```
! CAP372 - exercise 5 - 2019-09-15 - Parallel version
! Based on "Introduction to Parallel Programming: Matrix Multiplication"
! by Gergel V.P. :
! <http://www.lac.inpe.br/~stephan/CAP-372/matrixmult_microsoft.pdf>
! - Each subtask i hold one row of matrix A and one column of matrix B .
! - The result is stored in one element of C .
! - After that, every subtask i, 0≤i<n, transmits its column of matrix B
!   to the subtask with the number (i + 1) mod n .
! - When the number of processors ("p") is less than the number of
!   basic subtasks ("n"), each processor would execute several
!   inner products of matrix A rows and matrix B columns.
! - Each aggregated basic subtask ( = the calculation of one row of C)
!   determines several rows of the result matrix C.
! module load intel_psx/2019
! mpiifort -o ilist05parallel list05parallel.f90
! mpirun -n 2 ./ilist05parallel

program list05parallel
  use MPI
  implicit none

  ! C= A * B matrix initialization
  integer, parameter :: ms=4                ! matrix size ms x ms
  integer, parameter :: ps=2                ! expected processes
  integer, parameter :: st=ms/ps            ! # of stripes
  double precision, dimension(ms, ms) :: A, B, C=0 ! matrix definition
  double precision, dimension(st, ms) :: AR=0 ! stripes: A Rows
  double precision, dimension(ms, st) :: BC=0 ! stripes: B Cols
  double precision, dimension(st, st) :: CS=0 ! stripes: C Stripe
  double precision, dimension(ms) :: AT=0, BT=0 ! temporary

  integer :: stripe_i, stripe_countA, stripe_countB ! counters
  integer :: mi, mj, mk, source, dest, i            ! indexes
  integer :: iteration, subtask                     ! looping
  double precision :: time1, time2                  ! time elapsed
  double precision :: temp1                          ! temporary value

  ! MPI Initialization
  integer :: my_rank, sender, p, ierr               ! mpi variables
  integer :: tagC=0                                 ! variable tag
  integer, parameter :: tag=0, tagB=99999           ! fixed tag
  integer, dimension(MPI_STATUS_SIZE) :: status     ! mpi status
  call MPI_Init(ierr)                               ! mpi initialize
  call MPI_Comm_rank(MPI_COMM_WORLD, my_rank, ierr) ! mpi processor's id
  call MPI_Comm_size(MPI_COMM_WORLD, p, ierr)       ! mpi # of processors
  if (p .ne. ps) then                               ! abort if ps <> p
    if (my_rank .eq. 0) then
      print*, "Error: # of mpi processes is", p, " and the expected is", ps
    endif
    stop
  end if

  ! *** Each subtask hold one row|column of matrix A|B and return C ***
  ! The rank 0 has the entire matrix A|B .
  ! Matrix initialization:
  if (my_rank .eq. 0) then
    ! Example 4 x 4 (to check if everything is ok):
    A = transpose(reshape([ 5, 2, 6, 1,
                           0, 6, 2, 0,
                           3, 8, 1, 4,
                           1, 8, 5, 6 ], shape(A)))
    B = transpose(reshape([ 7, 5, 8, 0,
                           1, 8, 2, 6,
                           9, 4, 3, 8,
                           5, 3, 7, 9 ], shape(B)))

    ! The result should be:
    ! 96 68 69 69
  end if
```

```

! 24 56 18 52
! 58 95 71 92
! 90 107 81 142
! Example ms x ms:
! A = reshape([(mi, mi=1, ms*ms)], shape(A))
! B = reshape([(mi, mi=ms*ms+1, ms*ms*2)], shape(A))
! Elapsed time calculation:
call cpu_time(time1)
end if

! ##### BEGIN SEND DATA TO THE FIRST INTERATION #####
! Each subtask is one processor.
! Row A # is the same for the same subtask #, and then it is only
! necessary send|recv one time.
if (my_rank .eq. 0) then ! Rank 0 has the entire matrix.
do dest = 1, ms / st - 1 ! Send to the subtask.
do stripe_countA = 1, st ! Send the stripe.
mi = dest * st + stripe_countA ! calculate stripe line
! Send row by row. Can be optimized to send a block.
AT = A(mi, :)
call MPI_Send(AT, ms, MPI_DOUBLE_PRECISION, &
dest, tag, MPI_COMM_WORLD, ierr)
end do ! stripe_countA
end do ! dest
do mi = 1, st ! Rank 0: copy A stripe to AR
AR(mi, :) = A(mi, :)
end do ! mi
end if

! And then each subtask i, 0<i<n, receives. The i=0 is not necessary.
if (my_rank .ne. 0) then ! Rank > 0 receive the stripe.
do stripe_countA = 1, st ! Recv the A stripe.
! For sure it can be optimized, instead of receiving line by line.
call MPI_Recv(AT, ms, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
MPI_ANY_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)
AR(stripe_countA, :) = AT
end do ! stripe_countA
end if ! my_rank .ne. 0

! In the first interaction, the subtasks should receive initial B row.
! In the next interactions, each subtask send to another.
! First time B row send:
if (my_rank .eq. 0) then ! Rank 0 has the entire matrix.
interation = 0 ! Iteration of the algorithm.
do subtask = 0, ms / st - 1 ! Each subtask hold one matrix row.
do stripe_countB = 0, st - 1 ! Count for p < n .
! stripe column calculation
stripe_i = st*mod(subtask+p-interation,p)+stripe_countB
mk = stripe_i + 1 ! Matrix B column index
if (subtask .eq. 0) then
BC(:, mk) = B(:, mk) ! Rank 0 only copy
else
BT = B(:, mk) ! Send to rank > 0
dest = subtask
call MPI_Send(BT, ms, MPI_DOUBLE_PRECISION, &
dest, tag, MPI_COMM_WORLD, ierr)
endif ! subtask
end do ! stripe_countB
end do ! subtask
end if ! my_rank

! Recv B row that was sent
if (my_rank .ne. 0) then
do stripe_countB = 0, st - 1 ! Count for p < n
call MPI_Recv(BT, ms, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
MPI_ANY_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)
BC(:, stripe_countB + 1) = BT
end do ! stripe_countB
end if ! my_rank
! ##### END SEND DATA FOR THE FIRST INTERATION #####

```

```

! ++++++ START INTERACTION ALGORITHM ++++++
do iteration = 0, ms / st - 1 ! Algorithm iteration.
! Every subtask receives its row|col of matrix A|B
subtask = my_rank ! Each subtask hold one matrix A|B row|col
! >>>>>>>> START STRIPE <<<<<<<<<
do stripe_countA = 0, st - 1 ! A count for p < n
do stripe_countB = 0, st - 1 ! count for p < n
! stripe column calculation
stripe_i = st * mod(subtask+p-iteration,p) + stripe_countB
! matrix A line calculation
mi = st * subtask + stripe_countA + 1
temp1 = 0 ! multiplication: A line x B column
do mj = 1, ms ! count A column
mk = stripe_i + 1 ! B column index
! C = A x B. Can be optimized using intrinsic functions
temp1 = temp1 + AR(stripe_countA+1, mj) * BC(mj, stripe_countB+1)
end do
CS(stripe_countA+1, stripe_countB+1) = temp1 ! store C element
end do
end do ! >>>>>>>> END STRIPE <<<<<<<<<

! $$$$$$$$ send C stripe to rank 0 $$$$$$$$
! Each subtask should return the calculated C to rank 0
if (my_rank .ne. 0) then ! Rank 0 does not need to send
tagC = iteration
do mj = 1, st ! Send one stripe.
call MPI_Send(CS(:,mj), st, MPI_DOUBLE_PRECISION, 0, tagC, &
MPI_COMM_WORLD, ierr)
end do
end if

! Rank 0 is responsible for completing the matrix C
if (my_rank .eq. 0) then
! For sure there is a way to optimize this part of code
! SEND C stripe
do stripe_countA = 0, st - 1
do stripe_countB = 0, st - 1
stripe_i = st * mod(0 + p - iteration, p) + stripe_countB + 1
mi = stripe_countA + 1
C(mi, stripe_i) = CS(stripe_countA + 1, stripe_countB + 1)
end do
end do
! RECV C stripe
do source = 1, ps - 1 ! recv from ranks > 0 and store in C
! Receive the corresponding C
do mj = 1, st ! Recv one stripe.
call MPI_Recv(CS(:, mj), st * st, MPI_DOUBLE_PRECISION, source, &
MPI_ANY_TAG, MPI_COMM_WORLD, status, ierr)
end do
sender = status(MPI_SOURCE)
i = status(MPI_TAG)
! For sure there is a way to optimize this part of code
do stripe_countA = 0, st - 1
do stripe_countB = 0, st - 1
mi = st * sender + stripe_countA + 1 ! Calculate C line index
stripe_i = st * mod(sender + p - i, p) + stripe_countB
mk = stripe_i + 1 ! B column index
C(mi, mk) = CS(stripe_countA + 1, stripe_countB + 1)
end do ! stripe_countA
end do ! stripe_countB
end do ! source
end if ! $$$$$$$$ END OF SEND C $$$$$$$$

! @@@@@@@@@@ send and recv B @@@@@@@@@@
! At the end of one iteration, every subtask transmits its column
! of matrix B to the subtask with the number (i + 1) mod n .
! NOTE: the count index in Intel MPI implementation is a 4-byte integer,
! so the maximum allowed value is 2^31-1 : 2.147.483.647 .
! I had a lot o trouble with this.....
if (iteration .lt. (ms / st - 1)) then ! Not need in the last iteration.
dest = mod(my_rank + 1, ps)
do mj = 1, st ! Send one stripe.

```

```

! tagAB ensures no confusion with matrix C
call MPI_Send(BC(:,mj), ms, MPI_DOUBLE_PRECISION, dest, &
              tagB, MPI_COMM_WORLD, ierr)
end do
! Each subtask recv
do mj = 1, st ! Recv one stripe.
! tagAB ensures no confusion with matrix C
call MPI_Recv(BC(:,mj), ms, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
              tagB, MPI_COMM_WORLD, status, ierr)
end do
end if

end do ! ++++++ END INTERATION ++++++

! %%%%%%%%% SHOW THE RESULT %%%%%%%%%
if (my_rank .eq. 0) then
! Uncomment to show the multiplication result C
print*, "Result Matrix C ="
do mi = 1, ms
print*, int(C(mi,:))
end do
call cpu_time(time2) ! elapsed time calculation
write(*, "(A, I5, A, I3, A, F8.4)") "Matrix size:", ms, &
", Processors:", p, ", Elapsed time [s]:", time2 - time1
end if

call MPI_Finalize(ierr)
end program list05parallel

```

Rodando a matriz de teste para verificar o funcionamento:

```

$ mpif90 -Og -Wall -fcheck=all list05parallel.f90
$ mpirun -n 2 ./a.out
Result Matrix C =
      96      68      69      69
      24      56      18      52
      58      95      71      92
      90     107      81     142
Matrix size:   4, Processors: 2, Elapsed time [s]: 0.0001

```

O programa foi feito de forma rápida e pode ser melhorado. O foco foi fazer funcionar em uma semana. Destes 7 dias, 1 dia foi gasto com Fortran e algoritmo. Os demais 6 dias foram gastos com o MPI, sendo 1 dia só com um problema intermitente com um MPI\_Send que em alguns casos travava silenciosamente e não emitia nenhuma mensagem....

A ideia do programa era fazer um passo-a-passo tentando seguir o que está escrito nos slides. Ele ainda não está finalizado. É apenas uma primeira versão para atender o prazo de entrega. O programa começa definindo variáveis, inicializando o MPI, e construindo as matrizes. Em seguida envia os dados para a primeira interação do algoritmo. Na sequência executa a primeira interação. Depois envia o resultado parcial para o rank 0 colocar na matriz C. Na sequência envia o “stripe” da matriz B para o próximo rank conforme descrição do algoritmo. O próximo passo é repetir a interação até que todos os elementos de C sejam calculados. O resultado final aparece na matriz C.

## Testes na máquina local (Intel i7) usando ifort :

```
$ mpiifort -g -check all -fpe0 -warn -traceback -debug extended -o ilist05parallel-512-2 list05parallel.f90
$ mpirun -n 2 ./ilist05parallel-512-2
Matrix size: 512, Processors: 2, Elapsed time [s]: 0.5320
```

```
$ mpiifort -g -check all -fpe0 -warn -traceback -debug extended list05parallel.f90
-o ilist05parallel-512-4
$ mpirun -n 4 ./ilist05parallel-512-4
Matrix size: 512, Processors: 4, Elapsed time [s]: 0.2669
```

```
$ mpiifort -g -check all -fpe0 -warn -traceback -debug extended list05parallel.f90
-o ilist05parallel-512-8
$ mpirun -n 8 ./ilist05parallel-512-8
Matrix size: 512, Processors: 8, Elapsed time [s]: 0.1552
```

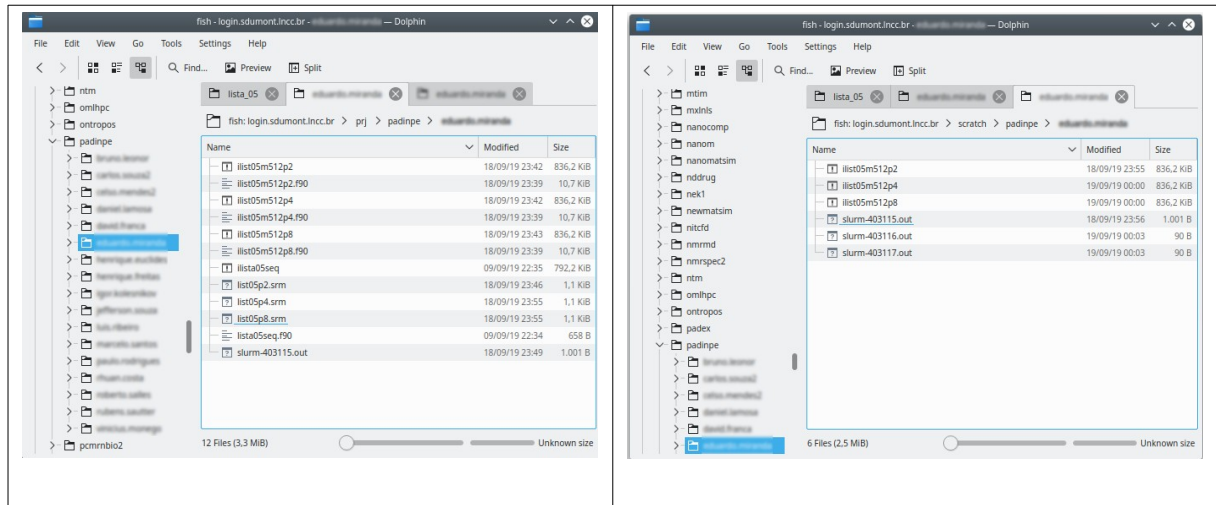
```
$ mpiifort -g -check all -fpe0 -warn -traceback -debug extended list05parallel.f90
-o ilist05parallel-1024-2
$ mpirun -n 2 ./ilist05parallel-1024-2
Matrix size: 1024, Processors: 2, Elapsed time [s]: 9.8530
```

```
$ mpiifort -g -check all -fpe0 -warn -traceback -debug extended list05parallel.f90
-o ilist05parallel-1024-4
$ mpirun -n 4 ./ilist05parallel-1024-4
Matrix size: 1024, Processors: 4, Elapsed time [s]: 2.3534
```

```
$ mpiifort -g -check all -fpe0 -warn -traceback -debug extended list05parallel.f90
-o ilist05parallel-1024-8
$ mpirun -n 8 ./ilist05parallel-1024-8
Matrix size: 1024, Processors: 8, Elapsed time [s]: 1.1655
```

# RODANDO NO SANTOS DUMONT

## Arquivos



### slurm-403115.out

```
dumont1292
sdumont1292
  linux-vdso.so.1 => (0x00007fff87daa000)
  libmpifort.so.12 =>
/opt/intel/parallel_studio_xe_2019/compilers_and_libraries_2019.3.199/linux/mpi/
intel64/lib/libmpifort.so.12 (0x00002b47c0382000)
  libmpi.so.12 =>
/opt/intel/parallel_studio_xe_2019/compilers_and_libraries_2019.3.199/linux/mpi/
intel64/lib/release/libmpi.so.12 (0x00002b47c0740000)
  libdl.so.2 => /usr/lib64/libdl.so.2 (0x00002b47c18ce000)
  librt.so.1 => /usr/lib64/librt.so.1 (0x00002b47c1ad2000)
  libpthread.so.0 => /usr/lib64/libpthread.so.0 (0x00002b47c1cda000)
  libm.so.6 => /usr/lib64/libm.so.6 (0x00002b47c1ef6000)
  libc.so.6 => /usr/lib64/libc.so.6 (0x00002b47c21f8000)
  libgcc_s.so.1 => /usr/lib64/libgcc_s.so.1 (0x00002b47c25c5000)
  libfabric.so.1 =>
/opt/intel/parallel_studio_xe_2019/compilers_and_libraries_2019.3.199/linux/mpi/
intel64/libfabric/lib/libfabric.so.1 (0x00002b47c27db000)
  /lib64/ld-linux-x86-64.so.2 (0x00002b47c015e000)
Matrix size: 512, Processors: 2, Elapsed time [s]: 0.6700
```

### slurm-403116.out

```
sdumont1292
sdumont1292
Matrix size: 512, Processors: 4, Elapsed time [s]: 0.3713
```

### slurm-403117.out

```
sdumont1292
sdumont1292
Matrix size: 512, Processors: 8, Elapsed time [s]: 0.2085
```



# TEMPORIZAÇÃO DO PROCESSAMENTO REFERENTE AOS CÁLCULOS E ÀS COMUNICAÇÕES

Foi utilizado o seguinte programa, que é o mesmo anterior acrescentado de “MPI\_Wtime”:

```

program list05p
! CAP372 - exercise 5 - 2019-09-15 - Parallel version
! Based on "Introduction to Parallel Programming: Matrix Multiplication"
! by Gergel V.P. :
! <http://www.lac.inpe.br/~stephan/CAP-372/matrixmult_microsoft.pdf>
! - Each subtask i hold one row of matrix A and one column of matrix B .
! - The result is stored in one element of C .
! - After that, every subtask i, 0≤i<n, transmits its column of matrix B
!   to the subtask with the number (i + 1) mod n .
! - When the number of processors ("p") is less than the number of
!   basic subtasks ("n"), each processor would execute several
!   inner products of matrix A rows and matrix B columns.
! - Each aggregated basic subtask ( = the calculation of one row of C)
!   determines several rows of the result matrix C.
! module load intel_psxe/2019
! mpiifort -g -check all -fpe0 -warn -traceback -debug extended &
!           -o ilist05p ilist05p.f90
! mpirun -n 2 ./ilist05p
! or if using gnu: mpif90 -Og -fcheck=all list05p.f90

use MPI
implicit none

integer, parameter :: ms=512                ! matrix size ms x ms
integer, parameter :: ps=2                 ! expected processes
integer, parameter :: st=ms/ps             ! # of stripes
double precision, dimension(ms, ms) :: A, B, C=0 ! matrix definition
double precision, dimension(st, ms) :: AR=0 ! stripes: A Rows
double precision, dimension(ms, st) :: BC=0 ! stripes: B Cols
double precision, dimension(st, st) :: CS=0 ! stripes: C Stripe
double precision, dimension(ms) :: AT=0, BT=0 ! temporary
integer :: stripe_i, stripe_countA, stripe_countB ! counters
integer :: mi, mj, mk, source, dest, i       ! indexes
integer :: interation, subtask               ! looping
double precision :: t1, t2, ttot, tcal=0, tcom=0 ! time elapsed
double precision :: temp1                   ! temporary value

! MPI Initialization
integer :: my_rank, sender, p, ierr          ! mpi variables
integer :: tagC=0                           ! variable tag
integer, parameter :: &
tag=0, & ! general use tag
tagB=99999, & ! tag for B matrix send
tagTcom=99998, & ! tag for communication time elapsed
tagTcal=99997 ! tag for calculation time elapsed
integer, dimension(MPI_STATUS_SIZE) :: status ! mpi status
call MPI_Init(ierr) ! mpi initialize
call MPI_Comm_rank(MPI_COMM_WORLD, my_rank, ierr) ! mpi processor's id
call MPI_Comm_size(MPI_COMM_WORLD, p, ierr) ! mpi # of processors
if (p .ne. ps) then ! abort if ps <> p
  if (my_rank .eq. 0) then
    print*, "Error: # of mpi processes is", p, " and the expected is", ps
  endif
  stop
end if
if (my_rank .eq. 0) then
  ttot = MPI_Wtime() ! Elapsed time calculation.
  ! *** Each subtask hold one row|column of matrix A|B and return C ***
  ! The rank 0 has the entire matrix A|B .
  A = reshape([(mi, mi=1, ms*ms)], shape(A))
  B = reshape([(mi, mi=ms*ms+1, ms*ms*2)], shape(B))
end if

!##### BEGIN SEND DATA TO THE FIRST INTERATION #####

```

```

! Each subtask is one processor.
! Row A # is the same for the same subtask #, and then it is only
! necessary send|recv one time.
if (my_rank .eq. 0) then          ! Rank 0 has the entire matrix.
t1 = MPI_Wtime()                  ! Start elapsed time calculation.
do dest = 1, ms / st - 1          ! Send to the subtask.
do stripe_countA = 1, st          ! Send the stripe.
mi = dest * st + stripe_countA    ! calculate stripe line
! Send row by row. Can be optimized to send a block.
AT = A(mi, :)
call MPI_Send(AT, ms, MPI_DOUBLE_PRECISION, &
dest, tag, MPI_COMM_WORLD, ierr)
end do ! stripe_countA
end do ! dest
do mi = 1, st                      ! Rank 0: copy A stripe to AR
AR(mi, :) = A(mi, :)
end do ! mi
end if
! And then each subtask i, 0<i<n, receives. The i=0 is not necessary.
if (my_rank .ne. 0) then          ! Rank > 0 receive the stripe.
do stripe_countA = 1, st          ! Recv the A stripe.
! For sure it can be optimized, instead of receiving line by line.
call MPI_Recv(AT, ms, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
MPI_ANY_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)
AR(stripe_countA, :) = AT
end do ! stripe_countA
t2 = MPI_Wtime() - t1             ! Elapsed time to send initial data
tcom = tcom + t2
end if ! my_rank .ne. 0

! In the first interaction, the subtasks should receive initial B row.
! In the next interactions, each subtask send to another.
! First time B row send:
if (my_rank .eq. 0) then          ! Rank 0 has the entire matrix.
interaction = 0                   ! Iteration of the algorithm.
do subtask = 0, ms / st - 1       ! Each subtask hold one matrix row.
do stripe_countB = 0, st - 1      ! Count for p < n .
! stripe column calculation
stripe_i = st * mod(subtask + p - interaction, p) + stripe_countB
mk = stripe_i + 1                 ! Matrix B column index
if (subtask .eq. 0) then
BC(:, mk) = B(:, mk)             ! Rank 0 only copy
else
BT = B(:, mk)                    ! Send to rank > 0
dest = subtask
call MPI_Send(BT, ms, MPI_DOUBLE_PRECISION, &
dest, tag, MPI_COMM_WORLD, ierr)
endif ! subtask
end do ! stripe_countB
end do ! subtask
end if ! my_rank

! Recv B row that was sent
if (my_rank .ne. 0) then
do stripe_countB = 0, st - 1      ! Count for p < n
call MPI_Recv(BT, ms, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
MPI_ANY_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)
BC(:, stripe_countB + 1) = BT
end do ! stripe_countB
end if ! my_rank

!+++++++ START INTERACTION ALGORITHM ++++++
do interaction = 0, ms / st - 1    ! Algorithm iteration.
t1 = MPI_Wtime()                  ! Start elapsed time calculation.
! Every subtask receives its row|col of matrix A|B
subtask = my_rank                 ! Each subtask hold one matrix A|B row|col
! >>>>>>>>> START STRIPE <<<<<<<<<
do stripe_countA = 0, st - 1      ! A count for p < n
do stripe_countB = 0, st - 1      ! count for p < n
! stripe column calculation
stripe_i = st * mod(subtask+p-interaction,p) + stripe_countB
! matrix A line calculation
mi = st * subtask + stripe_countA + 1
temp1 = 0                         ! multiplication: A line x B column

```

```

do mj = 1, ms ! count A column
mk = stripe_i + 1 ! B column index
! C = A x B. Can be optimized using intrinsic functions
temp1 = temp1 + AR(stripe_countA+1, mj) * BC(mj, stripe_countB+1)
end do
CS(stripe_countA+1, stripe_countB+1) = temp1 ! store C element
end do
! >>>>>>>>> END STRIPE <<<<<<<<<<
! Each rank has their partial time and must send to the rank 0
t2 = MPI_Wtime() - t1
if (my_rank .ne. 0) then ! rank 0 accumulate
call MPI_Send(t2, 1, MPI_DOUBLE_PRECISION, 0, &
tagTcal, MPI_COMM_WORLD, ierr)
else
tcal = tcal + t2 ! rank 0 own time
call MPI_Recv(t2, 1, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
tagTcal, MPI_COMM_WORLD, status, ierr)
tcal = tcal + t2 ! other rank time
endif ! End calculation elapsed time calculation

! $$$$$$$$$$ send C stripe to rank 0 $$$$$$$$$$
! Each subtask should return the calculated C to rank 0
t1 = MPI_Wtime() ! Start elapsed time calculation.
if (my_rank .ne. 0) then ! Rank 0 does not need to send
tagC = interaction
do mj = 1, st ! Send one stripe.
call MPI_Send(CS(:,mj), st, MPI_DOUBLE_PRECISION, 0, tagC, &
MPI_COMM_WORLD, ierr)
end do
end if
! Rank 0 is responsible for completing the matrix C
if (my_rank .eq. 0) then
! For sure there is a way to optimize this part of code
! SEND C stripe
do stripe_countA = 0, st - 1
do stripe_countB = 0, st - 1
stripe_i = st * mod(0 + p - interaction, p) + stripe_countB + 1
mi = stripe_countA + 1
C(mi, stripe_i) = CS(stripe_countA + 1, stripe_countB + 1)
end do
end do
! RECV C stripe
do source = 1, ps - 1 ! recv from ranks > 0 and store in C
! Receive the corresponding C
do mj = 1, st ! Recv one stripe.
call MPI_Recv(CS(:, mj), st * st, MPI_DOUBLE_PRECISION, source, &
MPI_ANY_TAG, MPI_COMM_WORLD, status, ierr)
end do
sender = status(MPI_SOURCE)
i = status(MPI_TAG)
! For sure there is a way to optimize this part of code
do stripe_countA = 0, st - 1
do stripe_countB = 0, st - 1
mi = st * sender + stripe_countA + 1 ! Calculate C line index
stripe_i = st * mod(sender + p - i, p) + stripe_countB
mk = stripe_i + 1 ! B column index
C(mi, mk) = CS(stripe_countA + 1, stripe_countB + 1)
end do ! stripe_countA
end do ! stripe_countB
end do ! source
end if ! $$$$$$$$$$ END OF SEND C $$$$$$$$$$

! @@@@@@@@@@ send and recv B @@@@@@@@@@
! At the end of one iteration, every subtask transmits its column
! of matrix B to the subtask with the number (i + 1) mod n .
! NOTE: the count index in Intel MPI implementation is a 4-byte integer,
! so the maximum allowed value is 2^31-1 : 2.147.483.647 .
! I had a lot o trouble with this.....
if (interaction .lt. (ms / st - 1)) then ! Not need in the last iteration.
dest = mod(my_rank + 1, ps)
do mj = 1, st ! Send one stripe.
! tagAB ensures no confusion with matrix C
call MPI_Send(BC(:,mj), ms, MPI_DOUBLE_PRECISION, dest, &
tagB, MPI_COMM_WORLD, ierr)

```

```

end do
! Each subtask recv
do mj = 1, st ! Recv one stripe.
! tagAB ensures no confusion with matrix C
call MPI_Recv(BC(:,mj), ms, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
tagB, MPI_COMM_WORLD, status, ierr)
end do
end if

! Each rank has their partial time and must send to the rank 0
t2 = MPI_Wtime() - t1
if (my_rank .ne. 0) then ! rank 0 accumulate
call MPI_Send(t2, 1, MPI_DOUBLE_PRECISION, 0, &
tagTcom, MPI_COMM_WORLD, ierr)
else
tcom = tcom + t2 ! rank 0 own time
call MPI_Recv(t2, 1, MPI_DOUBLE_PRECISION, MPI_ANY_SOURCE, &
tagTcom, MPI_COMM_WORLD, status, ierr)
tcom = tcom + t2 ! other rank time
endif ! End communication elapsed time calculation

end do ! ++++++ END INTERATION ++++++

!%%%%%%%%%% SHOW THE RESULT %%%%%%%%%%%
if (my_rank .eq. 0) then
ttot = MPI_Wtime() - ttot ! Elapsed time calculation.
write(*, '(a, i5, 3x, a, i2, 3x, a, f7.4, 3x, a, f7.4, 3x, a, f7.4)') &
'Matrix size:', ms, &
'Processors:', p, &
'Total time:', ttot, &
'Comm time:', tcom, &
'Calc time:', tcal
end if

call MPI_Finalize(ierr) ! mpi finalize
end program list05p

```

Resultado obtido rodando na máquina local:

```

$ mpif90 -Og -fcheck=all list05p.f90
$ mpirun -n 2 ./a.out
Matrix size: 512 Processors: 2 Total time: 0.2387 Comm time: 0.0058 Calc time: 0.4517

```

onde:

Total time	-	tempo medido do início ao final do Rank 0
Comm time	-	soma de todos os tempos de comunicação dos Ranks
Calc time	-	soma de todos os tempos de todos os Ranks (soma dos Ranks que estão rodando em paralelo)

## SPEED UP


O Speedup é o tempo de um programa série dividido pelo tempo do programa paralelo:

$$S = 1.3891 / 0.6700 = 2,07$$

A Eficiência é o Speedup dividido pelo nro. de processadores:

$$E = 2,07 / 2 = 1,04$$

Matrix size	Serial Algorithm	2 processors		4 processors		8 processors	
		Time	Seep UP	Time	Seep UP	Time	Seep UP
512	1,33891	0,67	2,00	0,3713	3,61	0,2085	6,42



## REFERÊNCIAS

- <http://www.lac.inpe.br/~stephan/CAP-372/>
- <https://stackoverflow.com>
- <https://annefou.github.io/Fortran/>
- <http://www.mathcs.emory.edu/~cheung/Courses/561/Syllabus/syl.html>
- [https://web.stanford.edu/class/me200c/tutorial\\_90/](https://web.stanford.edu/class/me200c/tutorial_90/)
- <http://userweb.eng.gla.ac.uk/peter.smart/com/fortran.htm>
- <https://software.intel.com/en-us/fortran-compilers>
- <http://fortranwiki.org>
- [https://events.prace-ri.eu/event/176/contributions/57/attachments/148/296/Advanced\\_MPI\\_I.pdf](https://events.prace-ri.eu/event/176/contributions/57/attachments/148/296/Advanced_MPI_I.pdf)
- [https://people.sc.fsu.edu/~jburkardt/f\\_src/mpi/matvec\\_mpi.f90](https://people.sc.fsu.edu/~jburkardt/f_src/mpi/matvec_mpi.f90)
- [https://www.dartmouth.edu/~rc/classes/intro\\_mpi/](https://www.dartmouth.edu/~rc/classes/intro_mpi/)
- <https://en.wikibooks.org/wiki/Fortran/>
- <http://people.ds.cam.ac.uk/nmm1/MPI>
- [https://events.prace-ri.eu/event/176/contributions/71/attachments/162/317/Advanced\\_MPI\\_I.pdf](https://events.prace-ri.eu/event/176/contributions/71/attachments/162/317/Advanced_MPI_I.pdf)